

CodeWhisperer Local – Your Private, Offline AI Coding Copilot

"Run. Break. Fix. Learn. All on your machine. Zero cloud. Zero APIs. Zero data leaks."

Project Title

CodeWhisperer Local

An intelligent, privacy-first, desktop-native coding assistant powered by Ollama (Mistral) – designed for developers who value control, speed, and security.

Core Vision

A sleek, local-first AI coding environment where you:

- **Write code** in a beautiful editor
- **Run it locally** – errors caught in a secure sandbox
- **Get AI-powered fixes** – not just syntax, but logic, environment, and dependency errors
- **Approve every fix** – full user control, no auto-overwrites
- **Explain code in 2 lines** – instant plain-English summaries
- **Refactor on demand** – cleaner, more efficient, Pythonic code

All 100% offline. All private. All yours.

Features (Small to Big – Nothing Missed)

1. Beautiful Local Code Editor

- **Monaco-based editor** (same as VS Code) with syntax highlighting, bracket matching, line numbers
- **Clean, distraction-free UI** – no clutter, no ads, no bloat
- **Persistent local storage** – your code stays on your machine

2. "Run" Button – Local Execution with Safety

- Runs your code in a **restricted, sandboxed environment** (no `exec()`, no global scope)
- Auto-captures **full runtime output, errors, and tracebacks**
- **Timeout protection** – kills runaway code after 10 seconds
- **Environment isolation** – runs in a clean Python subprocess to avoid polluting your system

3. AI-Powered Auto-Fix (The Heart of the Project)

➤ Error Detection & Context Capture

- Catches **syntax errors, runtime exceptions, logic bugs, and environment errors** (e.g., `ModuleNotFoundError`)
- Captures **full error type, message, traceback, and surrounding code context**
- Detects **missing packages** (`pip install pandas`) and suggests install commands
- Understands **environment context** – virtual env? Conda? Global? – and tailors fixes accordingly

➤ Smart Fix Generation via Ollama (Mistral)

- Sends structured error + code to **locally running Mistral via Ollama**
- Returns **minimal, precise fix** – no essays, no fluff
- Uses **specialized prompts** for different error types (syntax, runtime, import, logic)

➤ User-Controlled Fix Workflow

- **Never auto-applies fixes** – always shows suggestion first
- **"Apply & Run" button** – applies fix + auto-re-runs code
- **"Try Another Fix" button** – sends current code + "This fix failed. Try a different approach." → gets new fix
- **"Copy to Clipboard" button** – for manual pasting
- **Max 3 attempts** – prevents infinite loops. After 3 tries, shows final suggestion + manual help tips

➤ Environment & Dependency Intelligence

- If `import pandas` fails → suggests `pip install pandas`
- If `pip install` fails → suggests `pip install --upgrade pip` or `conda install pandas`
- Adds **pro tips**: "Activate your virtual env first" with OS-specific commands
- Never runs `pip install` without explicit user consent – **safety first**

4. "Explain" Button – 2-Line Code Summary

- Click → AI generates a **crisp, 2-line plain-English summary** of what the code *does* (not how)
- Perfect for legacy code, team projects, or your own forgotten scripts
- **Example:**
 Input: `def calculate_discount(price, discount_percent): return price * (1 - discount_percent / 100)`
 Output: Applies a percentage discount to a given price. Returns the final amount after discount.

5. "Refactor" Button – Code Beautification

- Click → AI returns a **cleaner, more efficient, Pythonic version** of your code
- Preserves functionality – no behavior change
- **Example:**
 Input: `result = []; for i in range(len(items)): if items[i] > 0: result.append(items[i] * 2)`
 Output: `result = [item * 2 for item in items if item > 0]`

6. Output & AI Panel – Clear, Organized Feedback

- **Output Panel:** Shows runtime results (green) or errors (red) – always visible
- **AI Fix Panel:** Appears only on error – clearly labeled "AI SUGGESTED FIX (Attempt X)"
- Shows **fix text + action buttons** – never mixes AI suggestions with your original code
- Includes **"Why This Fix?" tooltip** – hover for context (e.g., "This package is required for dataframes")

7. Safety & Control Features

- **No auto-execution** of fixes or `pip install` commands – user must click
- **No data leaves your machine** – all processing local
- **No cloud APIs** – zero cost, zero latency, zero privacy risk
- **Sandboxed execution** – protects your system from malicious or broken code

8. Extensibility & Future-Proofing

- **Modular design** – easy to add new features:
 - "Optimize Performance" button
 - "Add Type Hints" button
 - "Generate Unit Tests" button
- **Configurable** max attempts, timeout, and sandbox settings
- **Ready for multi-language support** (JS, SQL) – just swap the LLM prompt

Technology Stack

AI & Backend

- **LLM: Ollama (Mistral)** – runs locally, zero API cost, full privacy
- **Backend Framework: FastAPI** – lightweight, async, perfect for local tools
- **Sandbox:** Python `subprocess` + restricted environment – safe code execution
- **Error Parsing:** Custom Python logic to extract error type, message, traceback, and context

Frontend

- **Editor: Monaco Editor** – professional-grade, syntax-highlighted code editor
- **UI Framework:** Vanilla **HTML/CSS/JavaScript** – no React/Vue bloat, fast and lightweight
- **Styling:** Clean, modern CSS – responsive, accessible, beautiful

Local Infrastructure

- **Database:** None needed – all state stored in browser `localStorage` or temporary files
 - **Package Management:** Integrates with your local `pip` / `conda` – no bundled dependencies
 - **Deployment:** Runs 100% on localhost – no server required. Demo with `ngrok` if needed
-

Advantages (Why This Beats Everything Else)

For Developers

- **Privacy:** Your code never leaves your machine – no cloud, no APIs, no data leaks
- **Speed:** Local LLM + local execution = near-instant fixes (no network latency)
- **Control:** You approve every fix – no black-box auto-overwrites
- **Learning:** "Explain" and "Refactor" teach you better coding practices
- **Cost:** Zero API fees – Mistral runs free on your GPU/CPU

For Recruiters & Judges

- **Depth:** Shows mastery of local LLM orchestration, sandboxing, prompt engineering, and UI/UX
- **Originality:** No one in the dataset is building a user-approved, locally-run, agentic code debugger
- **Polish:** Beautiful frontend + professional UX – not a Streamlit/Gradio toy
- **Relevance:** Solves daily frustrations for every developer – not an academic exercise

For the Future

- **On-Device AI Trend:** Matches Apple Intelligence, Microsoft Copilot+ PCs — local, private, powerful
 - **Extensible:** Easy to add new features, languages, or integrations
 - **Deployable:** Package as a desktop app (Electron) or browser extension
-

Final Note

This isn't just another "AI code assistant."

This is **the future of intelligent, private, local development tools** — built by you, for you, on your machine.

No compromises. No cloud. No kidding.